

On the Information Bottleneck Theory of Deep Learning

Authors: A. M. Saxe, Y. Bansal, J. Dapello, M. Advani, A. Kolchinsky, B.D. Tracey, and D.D. Cox

Presenter: Zhongyuan Zhao

zhzhao@cse.unl.edu

Oct. 5th, 2018

A. M. Saxe, Y. Bansal, J. Dapello, M. Advani, A. Kolchinsky, B.D. Tracey, and D.D. Cox, “On the Information Bottleneck Theory of Deep Learning”, ICLR 2018, [Online] https://openreview.net/forum?id=ry_WPG-A-

Outline

- Part 1: Information Bottleneck Theory of Deep Learning
 - Information Theory Recap
 - Information Plane
 - IB Theory
- Part 2: Attacks on the IB theory
 - Counter Example
 - Summary
- Discussions



Part 1: Information Bottleneck Theory of Deep Neuron Networks

The Hebrew University of Jerusalem

Information Bottleneck Theory of DNN

- Who
 - Ravid Schwartz-Ziv, [Naftali Tishby](#), etc. The Hebrew University of Jerusalem
- What – an attempt to explain DNN
 - Training Dynamics
 - Learning Processes
 - Internal Representation
- How
 - Information Theory

- [1] R. Schwartz-Ziv and [N. Tishby](#). Opening the black box of deep neural networks via information. arXiv preprint arXiv:1703.00810, 2017
- [2] Michal Moshkovich and [Naftali Tishby](#). Mixing complexity and its applications to neural networks. 2017. URL <https://arxiv.org/abs/1703.00729>
- [3] [Naftali Tishby](#) and Noga Zaslavsky. Deep Learning and the information Bottleneck Principle. In Information Theory Workshop (ITW), 2015 IEEE, Pages 1-5. IEEE, 2015
- [4] [Naftali Tishby](#), Fernando C. Pereira, and William Bialek. The information bottleneck Method. In Proceedings of the 37-th Annual Allerton Conference on Communication, Control and Computing, 1999.

Information Theory: Entropy

A Discrete Random Variable X with possible values $\{x_1, \dots, x_n\}$

Information $I(x_i) = \log_2 \frac{1}{P(x_i)} = -\log_2 P(x_i) \geq 0$

Entropy $H(X) = E[I(X)] = \sum_{i=1}^n P(x_i)I(x_i) = -\sum_{i=1}^n P(x_i) \log_2 P(x_i)$

Joint Entropy $H(X, Y) = -\sum_{x \in \mathcal{X}} \sum_{y \in \mathcal{Y}} p(x, y) \log p(x, y).$

Conditional Entropy $H(X|Y) = -\sum_{i,j} p(x_i, y_j) \log \frac{p(x_i, y_j)}{p(y_j)}$

$$X : \{0, 1\}$$

$$P(0)=0.5, P(1)=0.5$$

$$I(0) = \log_2 \frac{1}{P(0)} = 1$$

$$H(X) = 0.5 \times 1 + 0.5 \times 1 = 1$$

$$P(0)=0.9, P(1)=0.1$$

$$I(0)=0.152, I(1) = 3.3219$$

$$H(X) = 0.469$$

Information Theory: Mutual Information

Mutual Information $I(X; Y) = \sum_{x \in X, y \in Y} p(x, y) \log\left(\frac{p(x, y)}{p(x)p(y)}\right) = H(X) - H(X|Y)$

Definition 5 *The mutual information $I(X; Y)$ measures how much (on average) the realization of random variable Y tells us about the realization of X , i.e., how by how much the entropy of X is reduced if we know the realization of Y .*

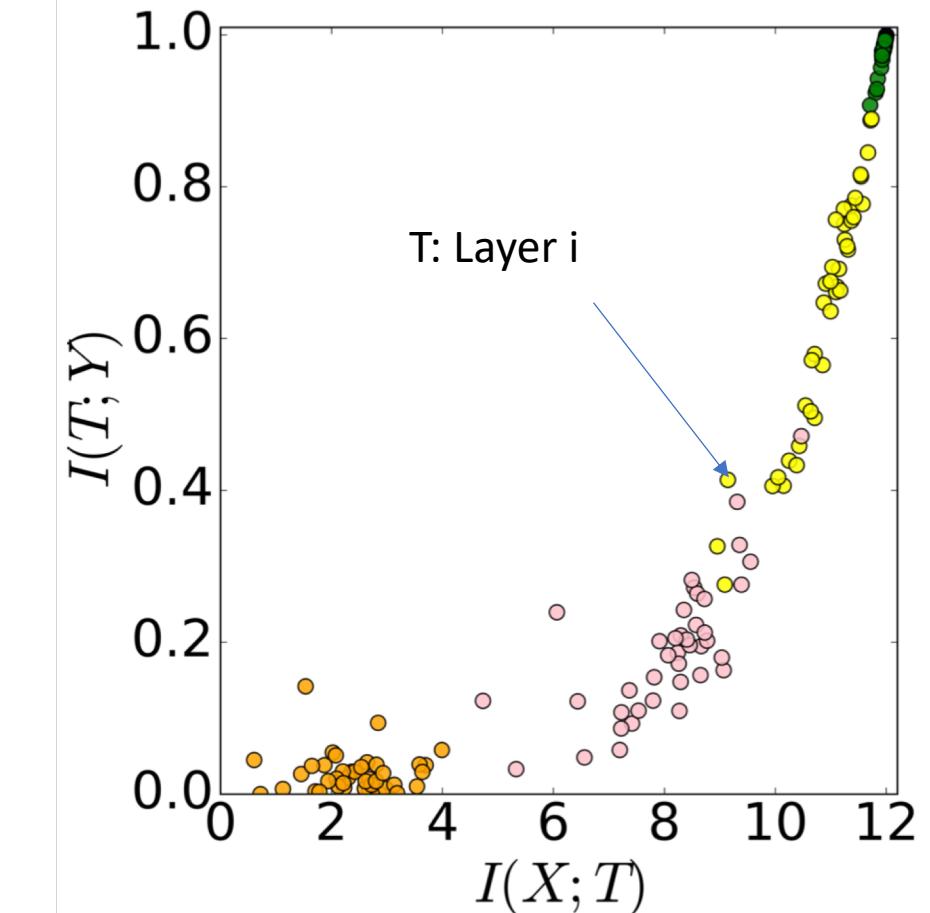
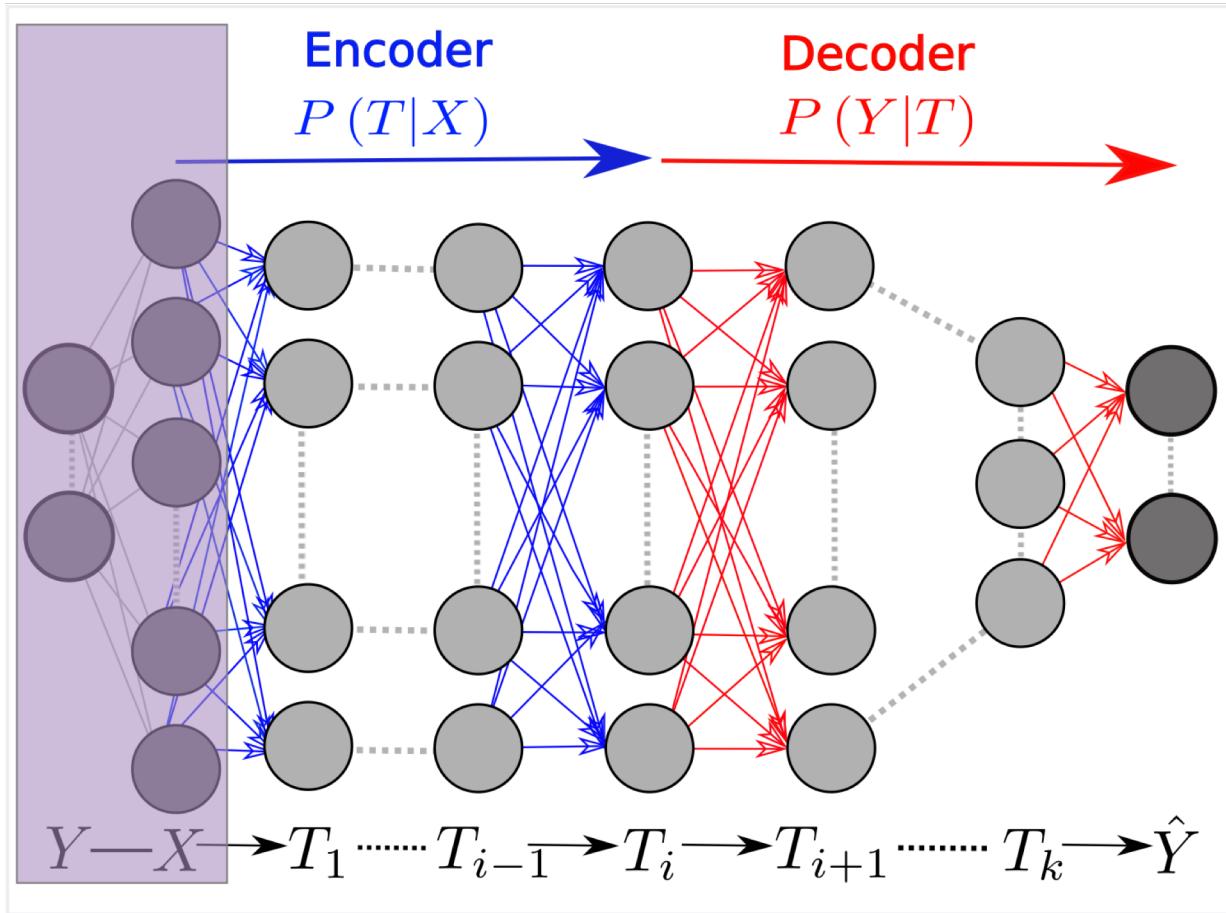
$$I(X; Y) = H(X) - H(X|Y)$$

Surprisingly, mutual information is symmetric; X tells us exactly as much about Y as Y tells us about X .

Theorem 2 *Symmetry of mutual information*

$$I(X; Y) = H(X) - H(X|Y) = H(Y) - H(Y|X) = I(Y, X).$$

Information Plane of DNN



Information Plane (Cont.)

- Invariance to Invertible Transformation

$$I(X; Y) = I(\psi(X); \phi(Y))$$

- Markov Chain

$$X \rightarrow Y \rightarrow Z$$

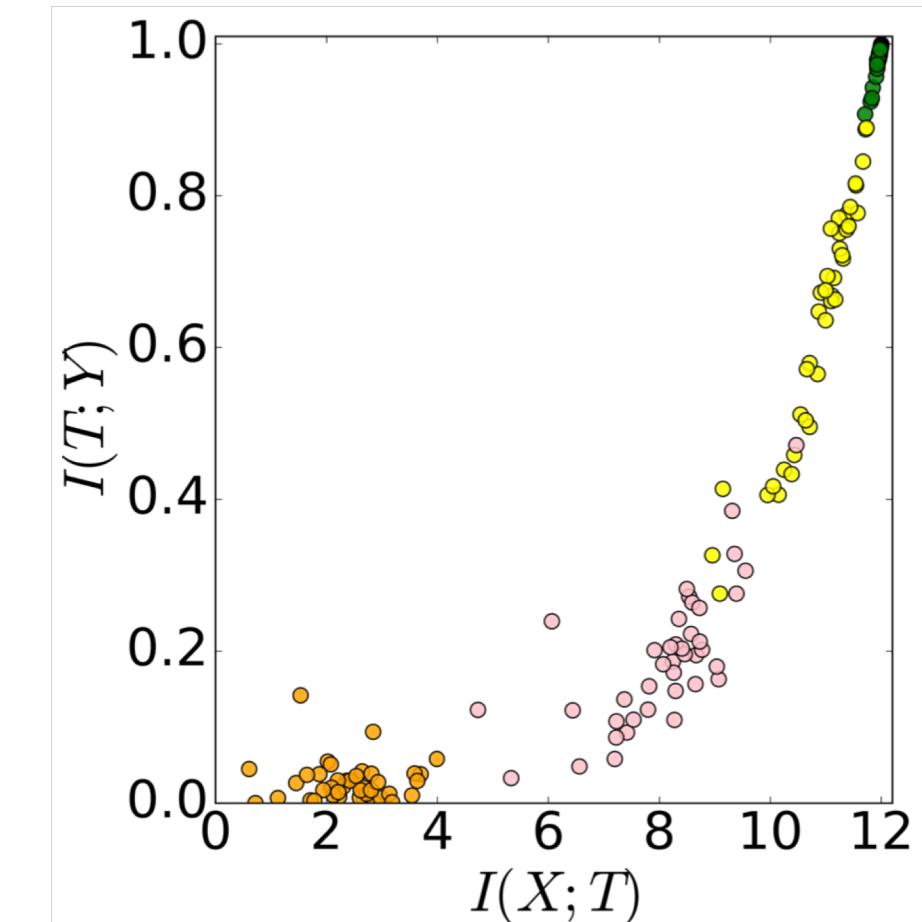
$$I(X; Y) \geq I(X; Z)$$

- Mutual Information Over Layers of DNN

$$\begin{aligned} I(X; Y) &\geq I(T_1; Y) \geq I(T_2; Y) \geq \dots \geq I(T_k; Y) \geq I(\hat{Y}; Y) \\ H(X) &\geq I(X; T_1) \geq I(X; T_2) \geq \dots \geq I(X; T_k) \geq I(X; \hat{Y}). \end{aligned}$$

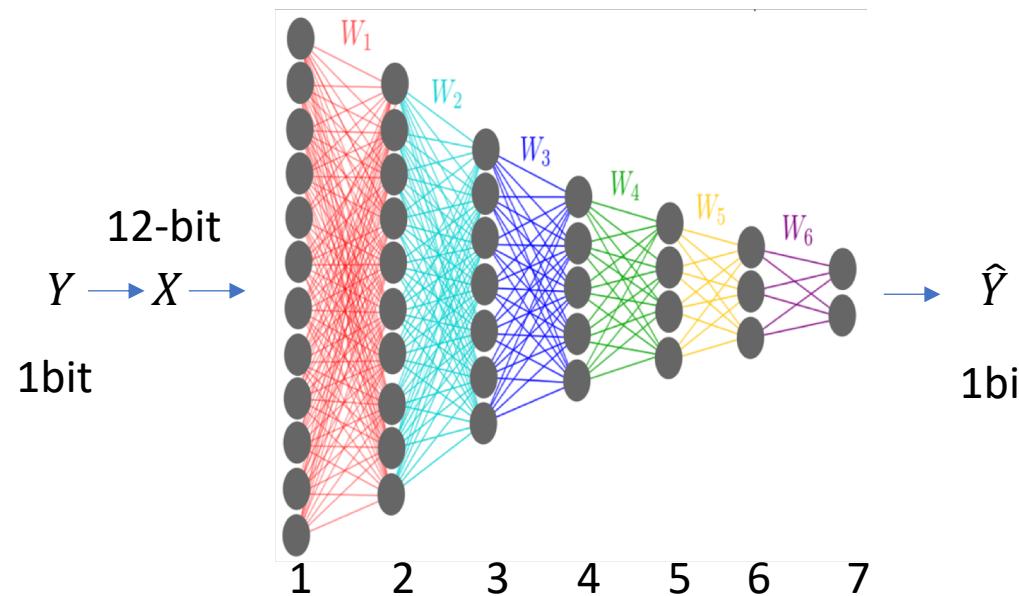
A **Markov chain** is "a [stochastic model](#) describing a [sequence](#) of possible events in which the probability of each event depends only on the state attained in the previous event".

--Wikipedia



Experimental Setup

- 7 Fully Connected Hidden Layers
 - 12-10-7-5-4-3-2
- Label Y : {0,1}
- Input X : 12-binary vector, total 4096, equally distributed



Mutual Information of Hidden Layers

- 30-bin for each neuron's output

$$t \in T_i$$

- Joint Distribution

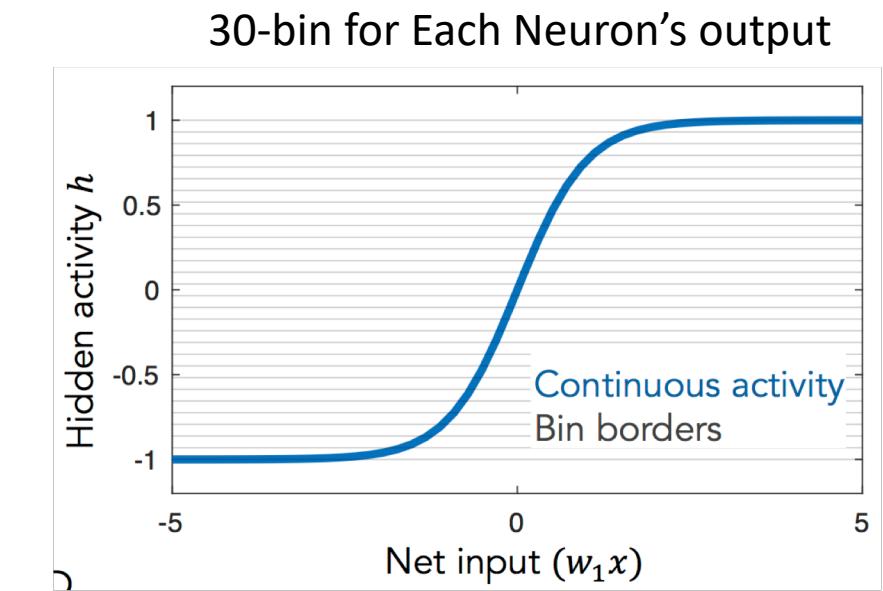
$$P(T_i, X)$$

$$P(T_i, Y) = \sum_x P(x, Y)P(T_i|x)$$

- Mutual Information

$$I(X; T_i)$$

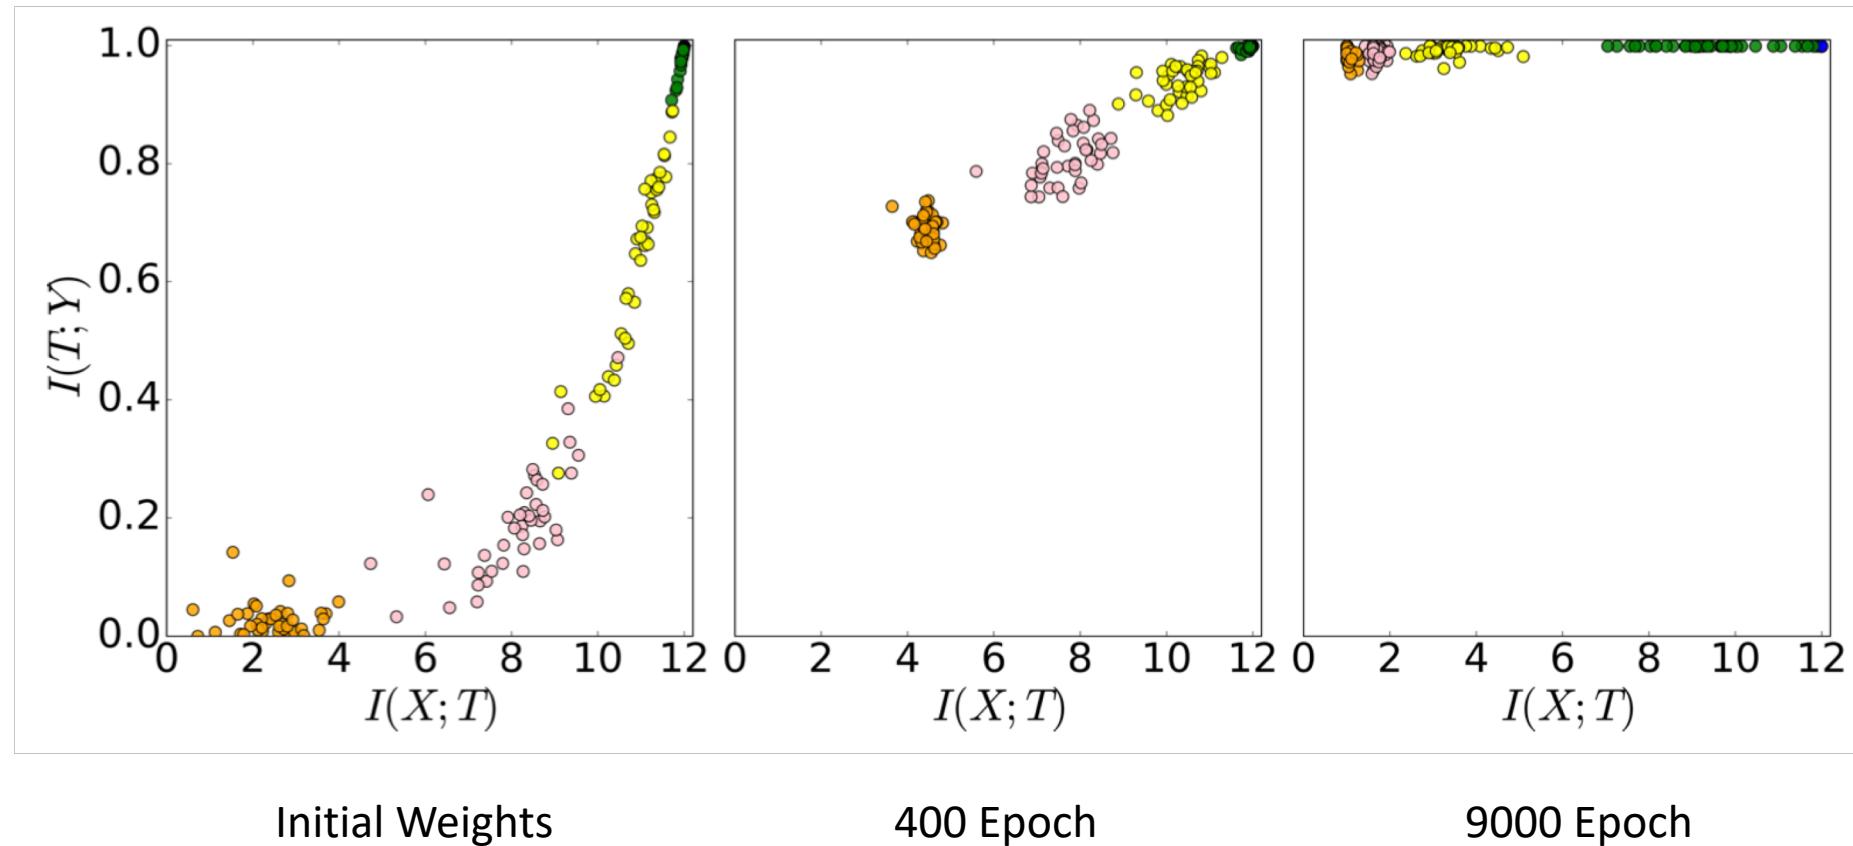
$$I(T_i; Y)$$



Training Process

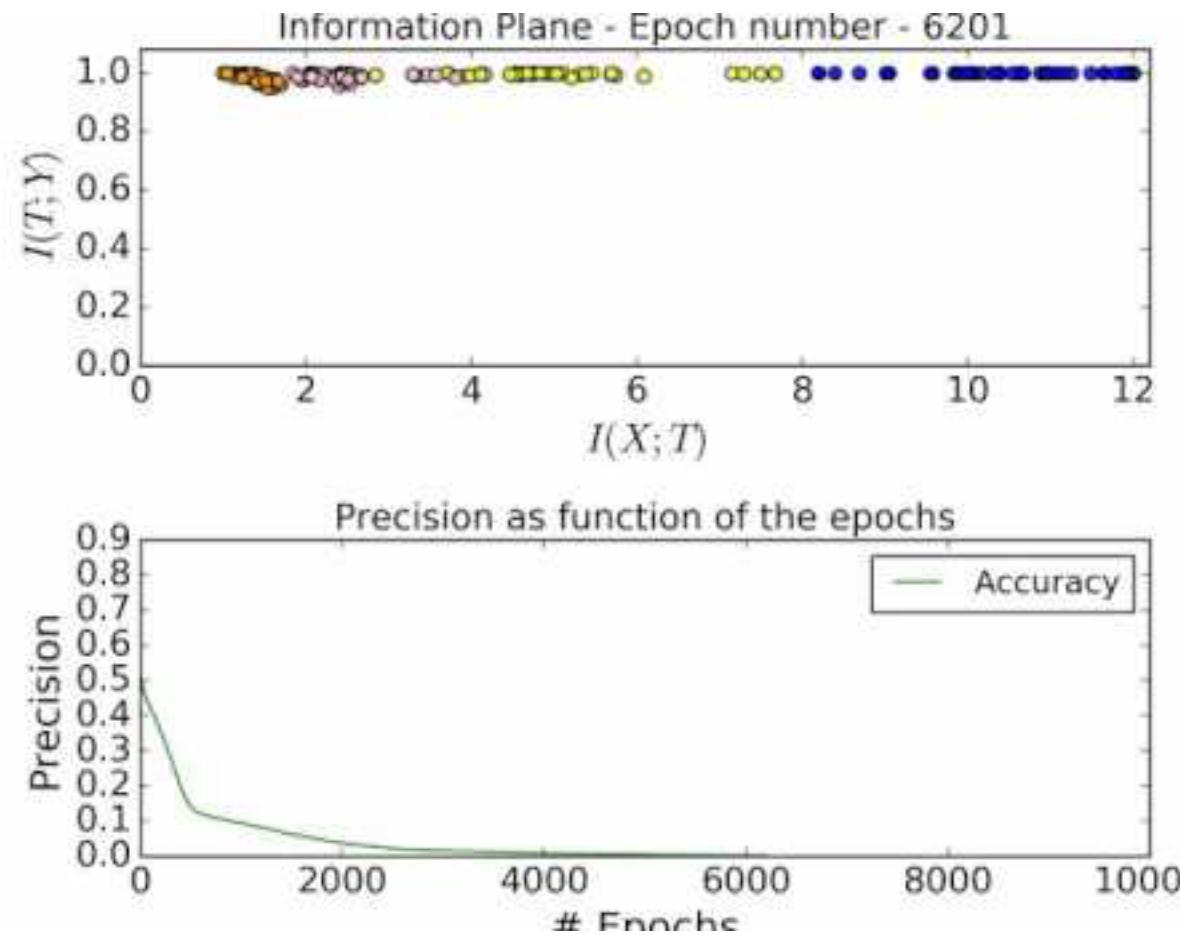
1. Color-coded hidden Layer
2. 50 randomized networks

<https://goo.gl/rygyIT>
<https://goo.gl/DQWuDD>



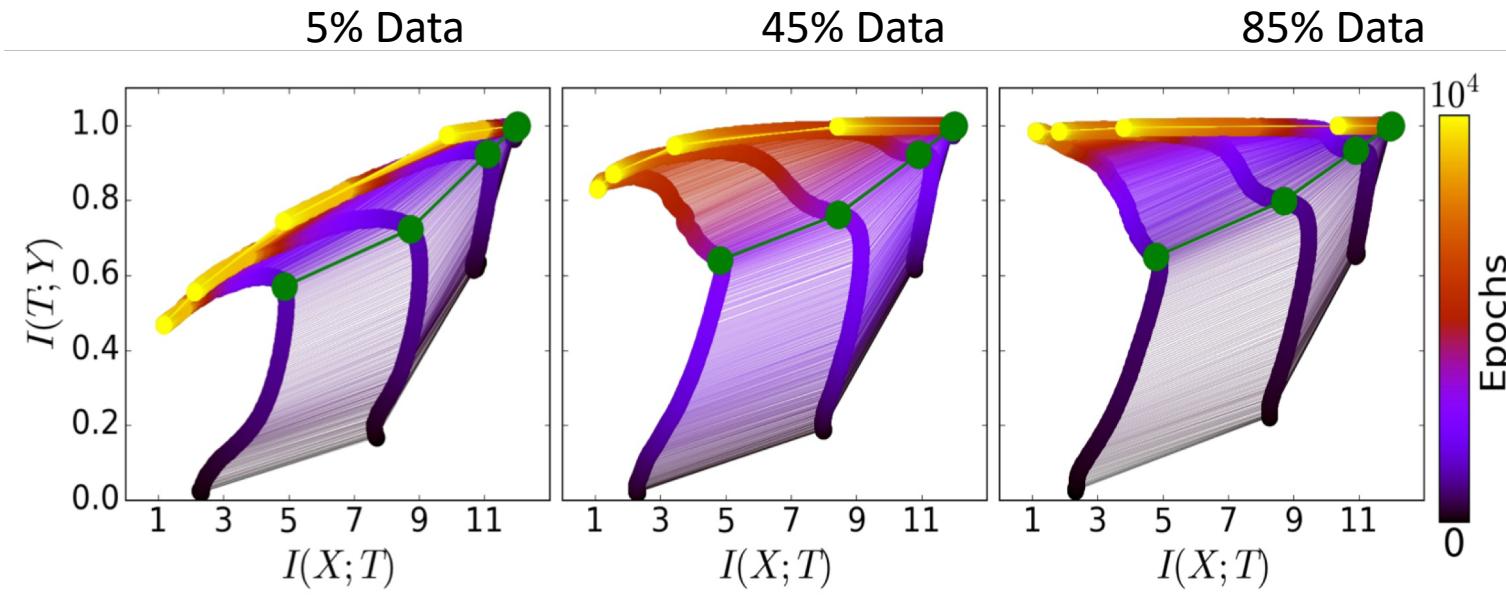
Training Process

- Fitting
- Compression



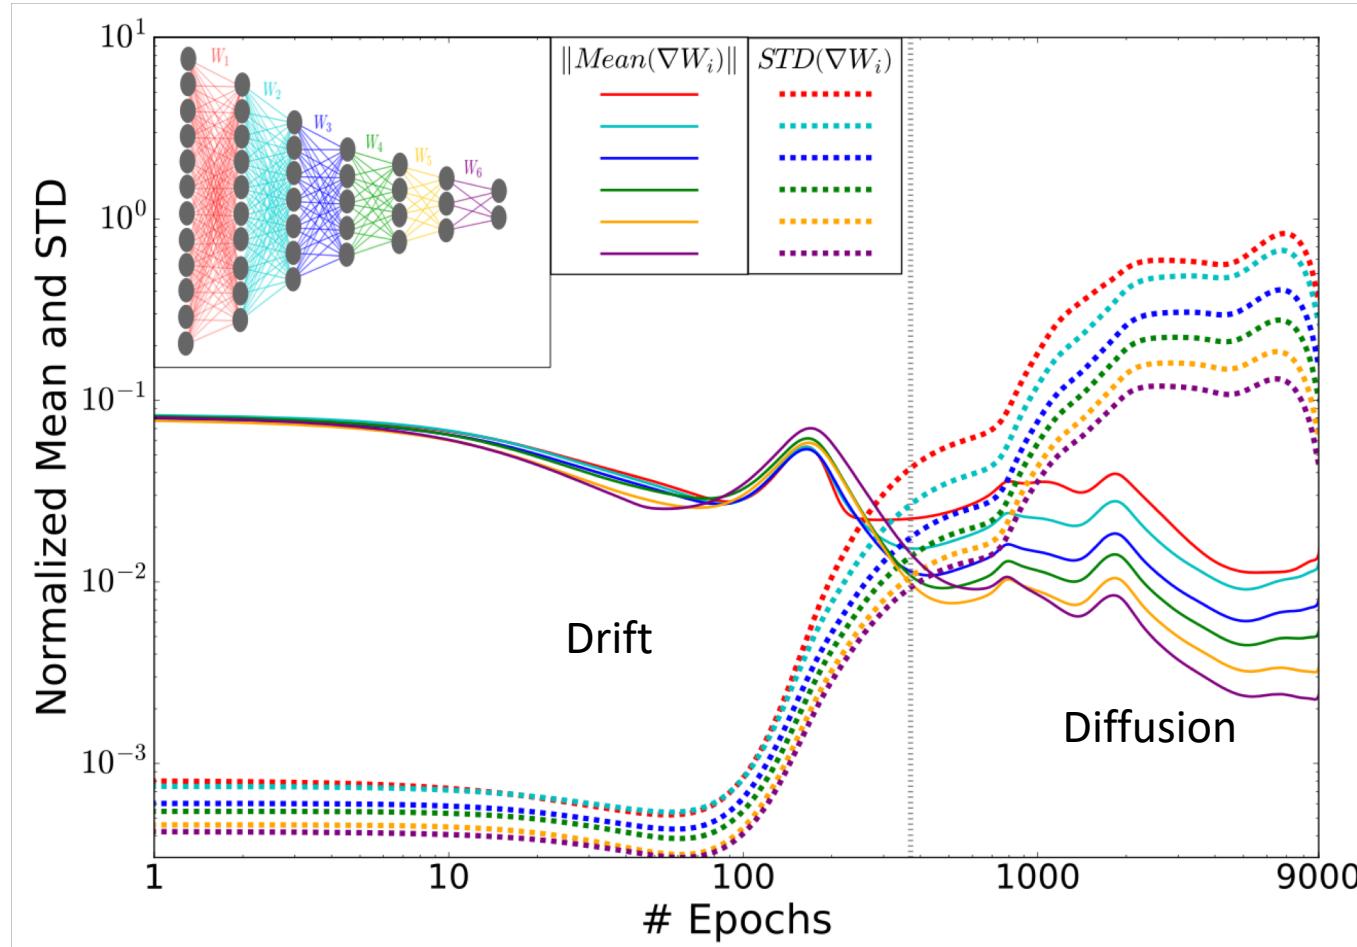
Two Phase Training Process

- Fitting (ERM)
 - Increase $I(X; T)$ Get relevant information from X
- Representation-Compression
 - Decrease $I(X; T)$ Drop Irrelevant Information from X



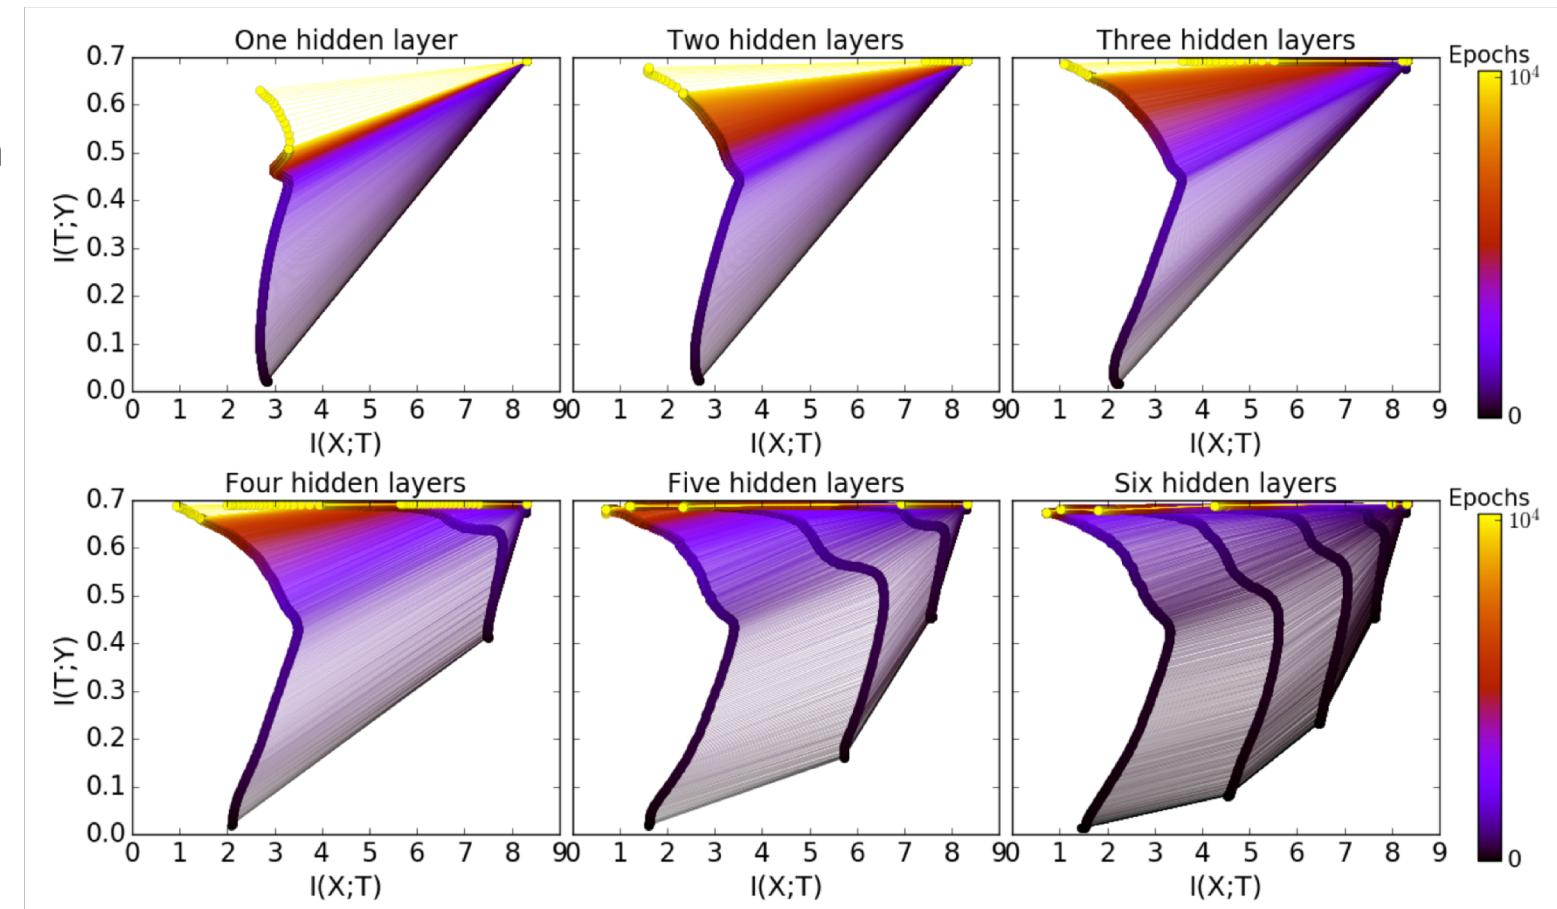
Drift and Diffusion Phases of SGD Optimization

- Weights' Stochastic Gradients

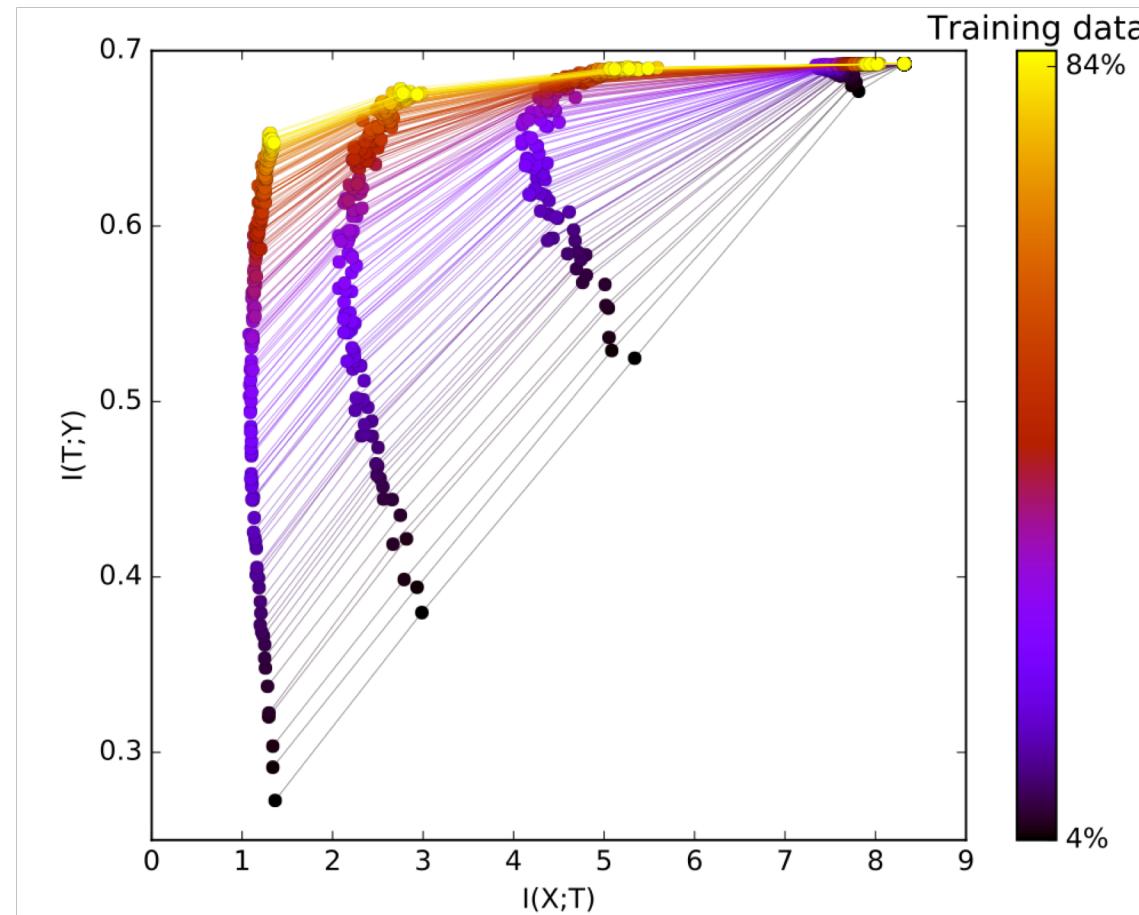


Computational Benefit of Hidden Layers

1. More Layers, less epoch
2. Compression phase is shorter if a layer start from a compressed layer
3. Compression is faster for deeper layers



Evolution of layers with training sample sizes



Converged Network with different training sample sizes

Summary of IB Theory on DNN

- Model
 - See DNN as a Markov Chain
 - Mutual Information of hidden layers with X, Y
- Learning Processes
 - Learning, Compression
- Training Dynamics
 - Drift, Diffusion
- Internal Representation
 - Information bottleneck set by training sample size
 - Hidden Layers brings computational benefit

Part 2: The Attack on IB Theory

Harvard University, Santa Fe Institute, MIT-IBM Watson AI Lab

Attacks on IB Theory on Deep Learning

- Who
 - A.M. Saxe, Y. Bansal, J. Dapello, M. Advani, (Harvard)
 - A. Kolchinsky, B.D. Tracey (Santa Fe Institute)
 - D.D. Cox (Harvard, MIT-IBM Watson AI Lab)
- What
 - Two Training Phases: Fitting, Compression
 - Compression phase related to generalization
 - Compression occurs due to diffusion-like behavior of SGD
- How
 - Counter Examples

Example 1: Relu v.s. tanh

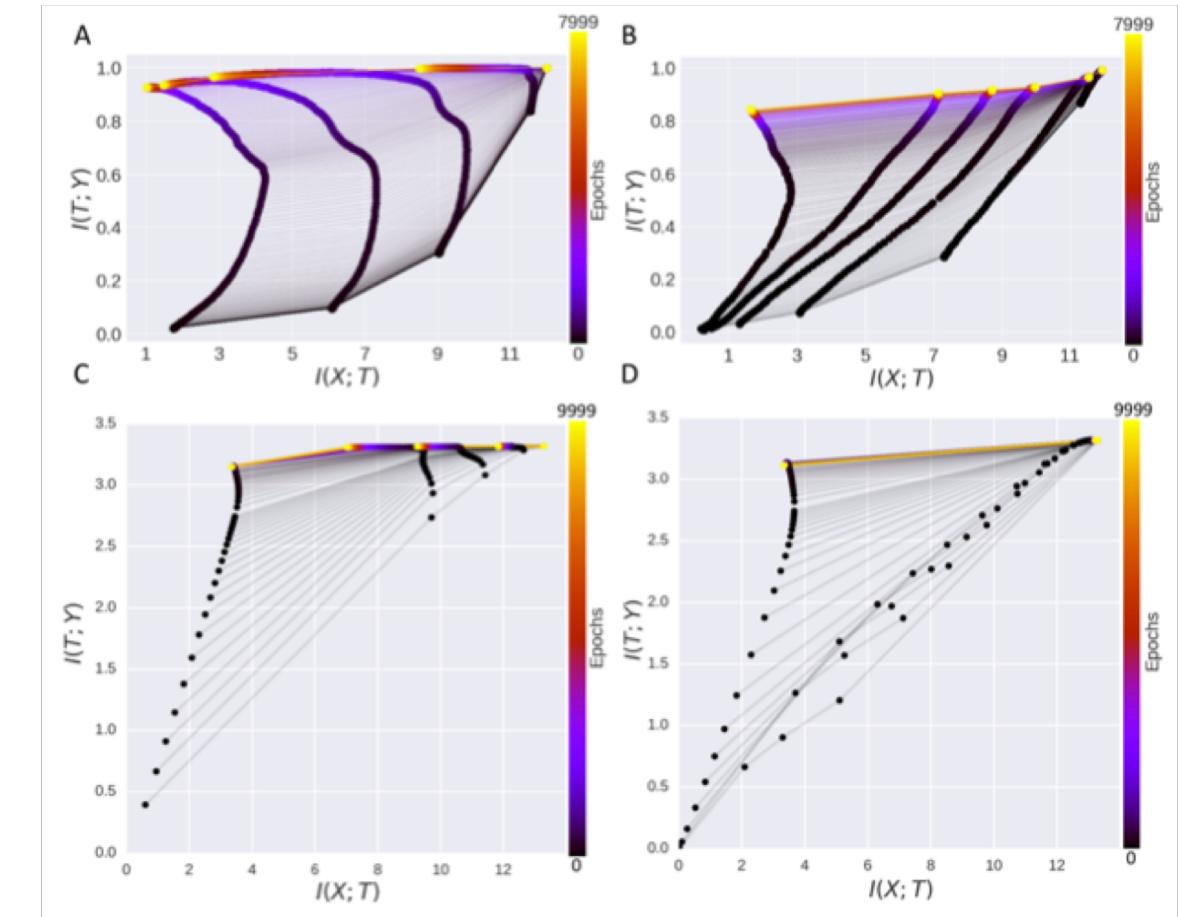
A: Replication of [1]

B: Replace tanh with Relu

X: 12-binary Vector
(12-10-7-5-4-3-2)

C: tanh network
D: Relu network

X: MNIST
(784-1024-20-20-10)



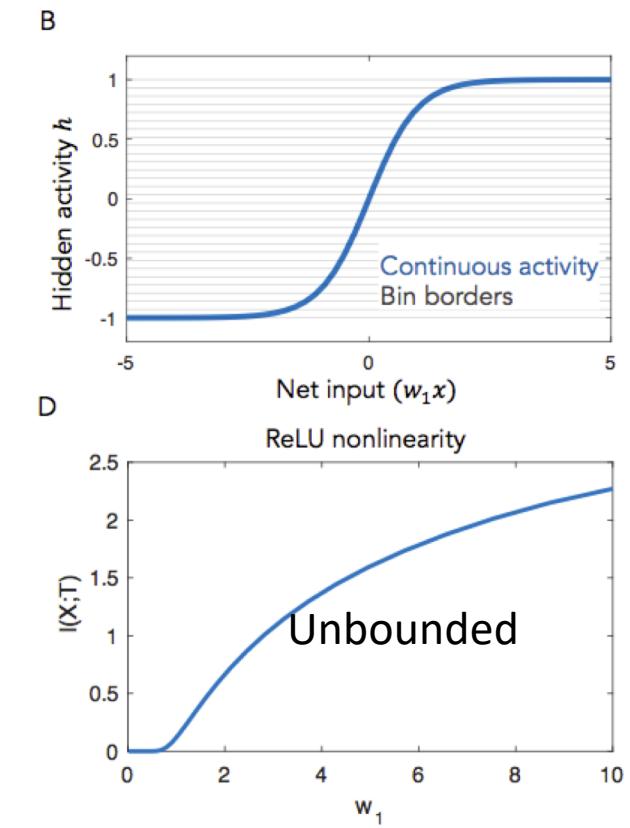
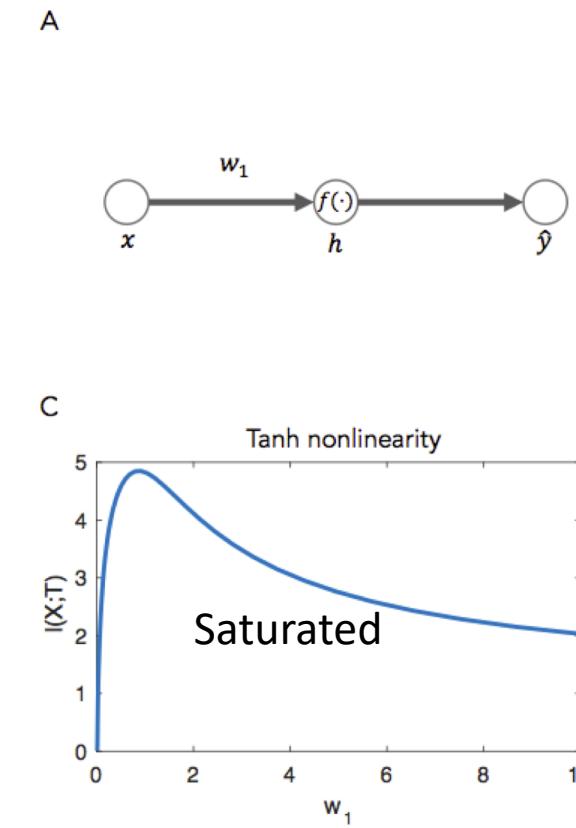
Minimal Model Illustration

1 hidden layer

Scalar Gaussian Input $X \sim N(0,1)$

$$\begin{aligned} I(T; X) &= H(T) - H(T|X) \\ &= H(T) \\ &= -\sum_{i=1}^N p_i \log p_i \end{aligned}$$

$$p_i = P(X \geq f^{-1}(b_i)/w_1 \text{ and } X < f^{-1}(b_{i+1})/w_1),$$



Example 2: Deep Linear Network

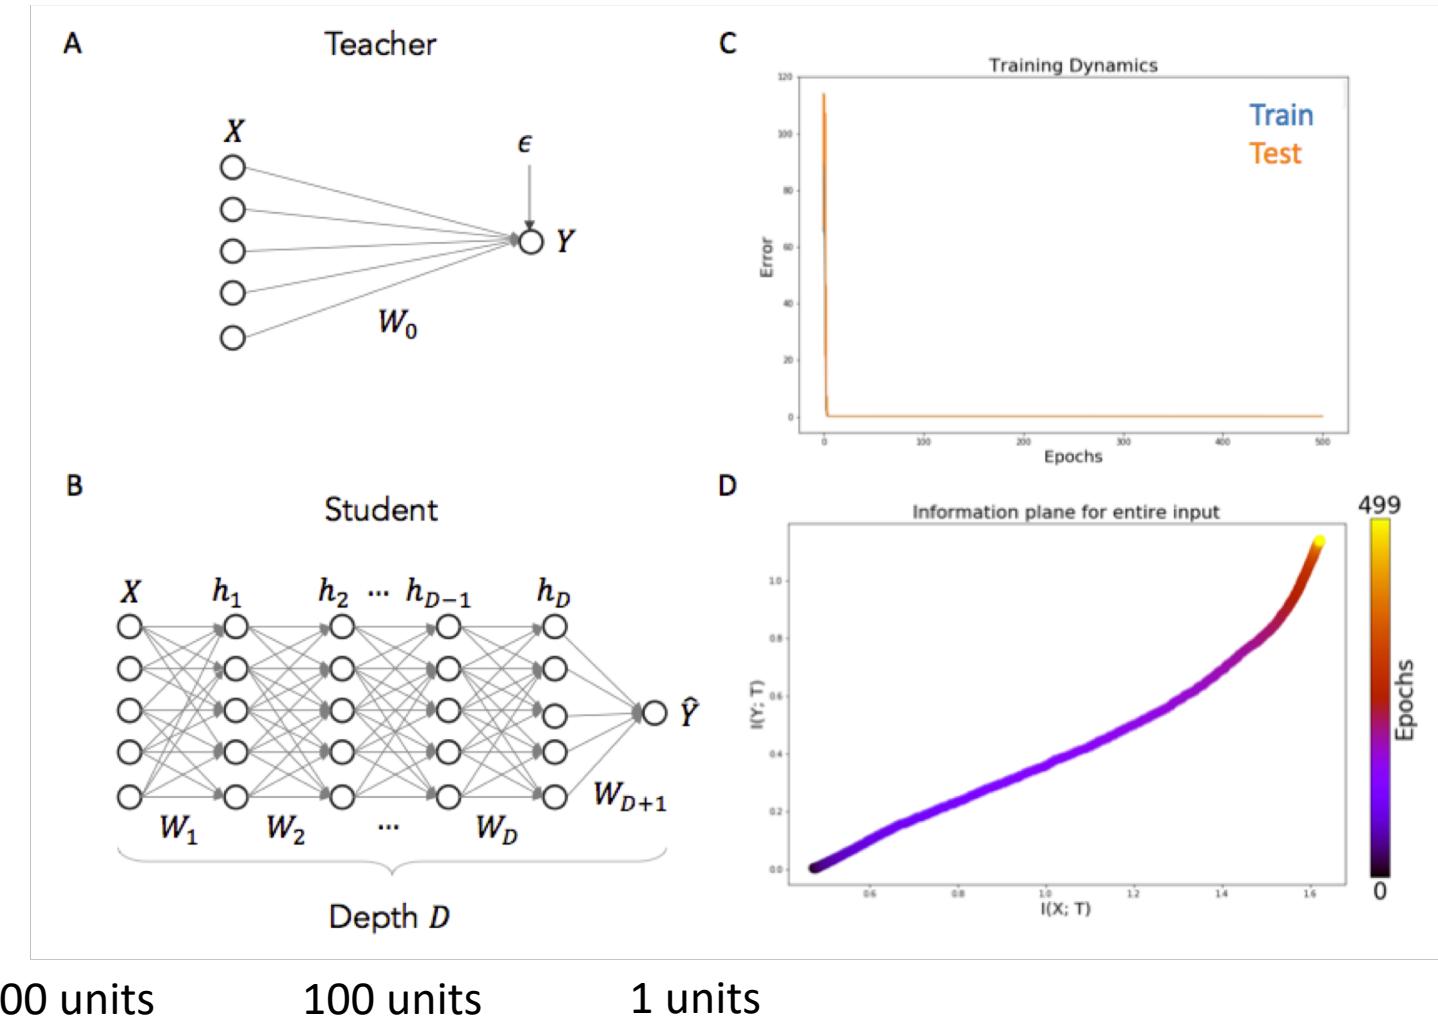
A: Linear Teacher Network

Gaussian input X

Add noise

B: Deep Linear Student Network

1 hidden layer

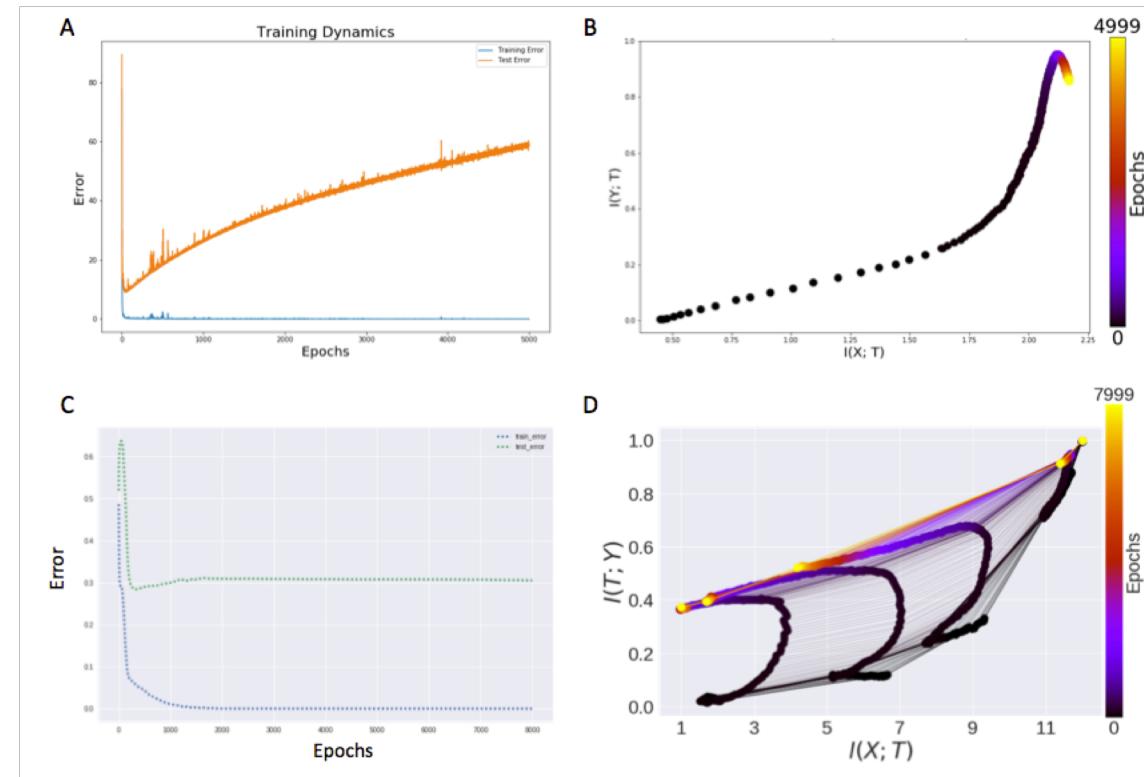


Example 3: Over-Training (over-fitting)

Deep Linear network

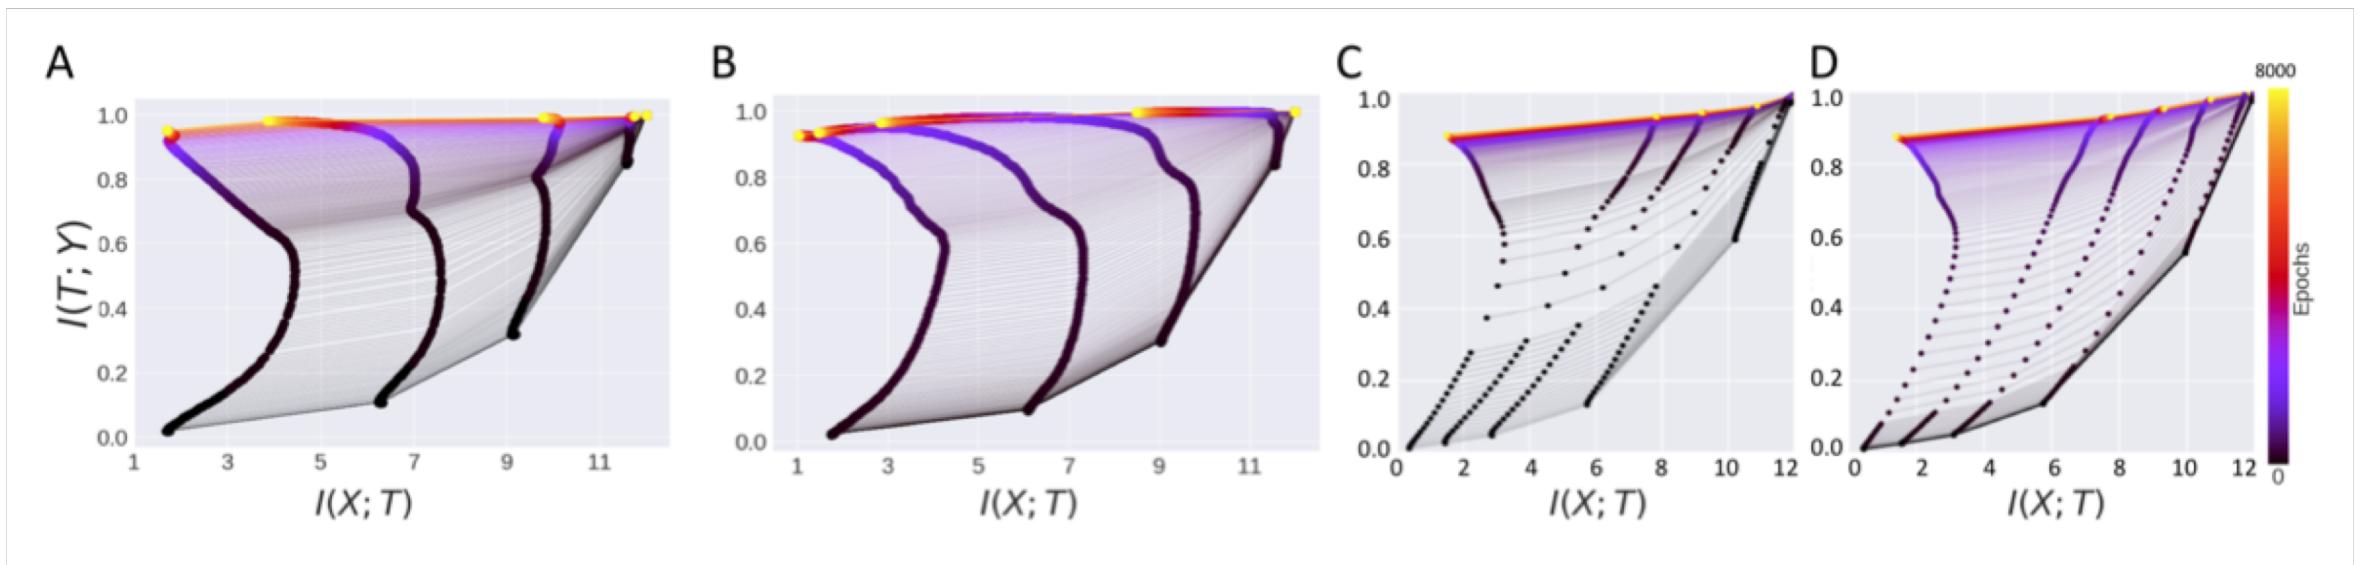
30% data

Tanh network



Cause of Compression: SGD v.s. BGD

- Stochastic Gradient Descent
 - random input samples, the weights evolve in a stochastic way during training
- Batch Gradient Descent
 - Full training dataset, no randomness or diffusion-like behavior in its updates



Tanh, SGD

Tanh, BGD

Relu, SGD

Relu, BGD

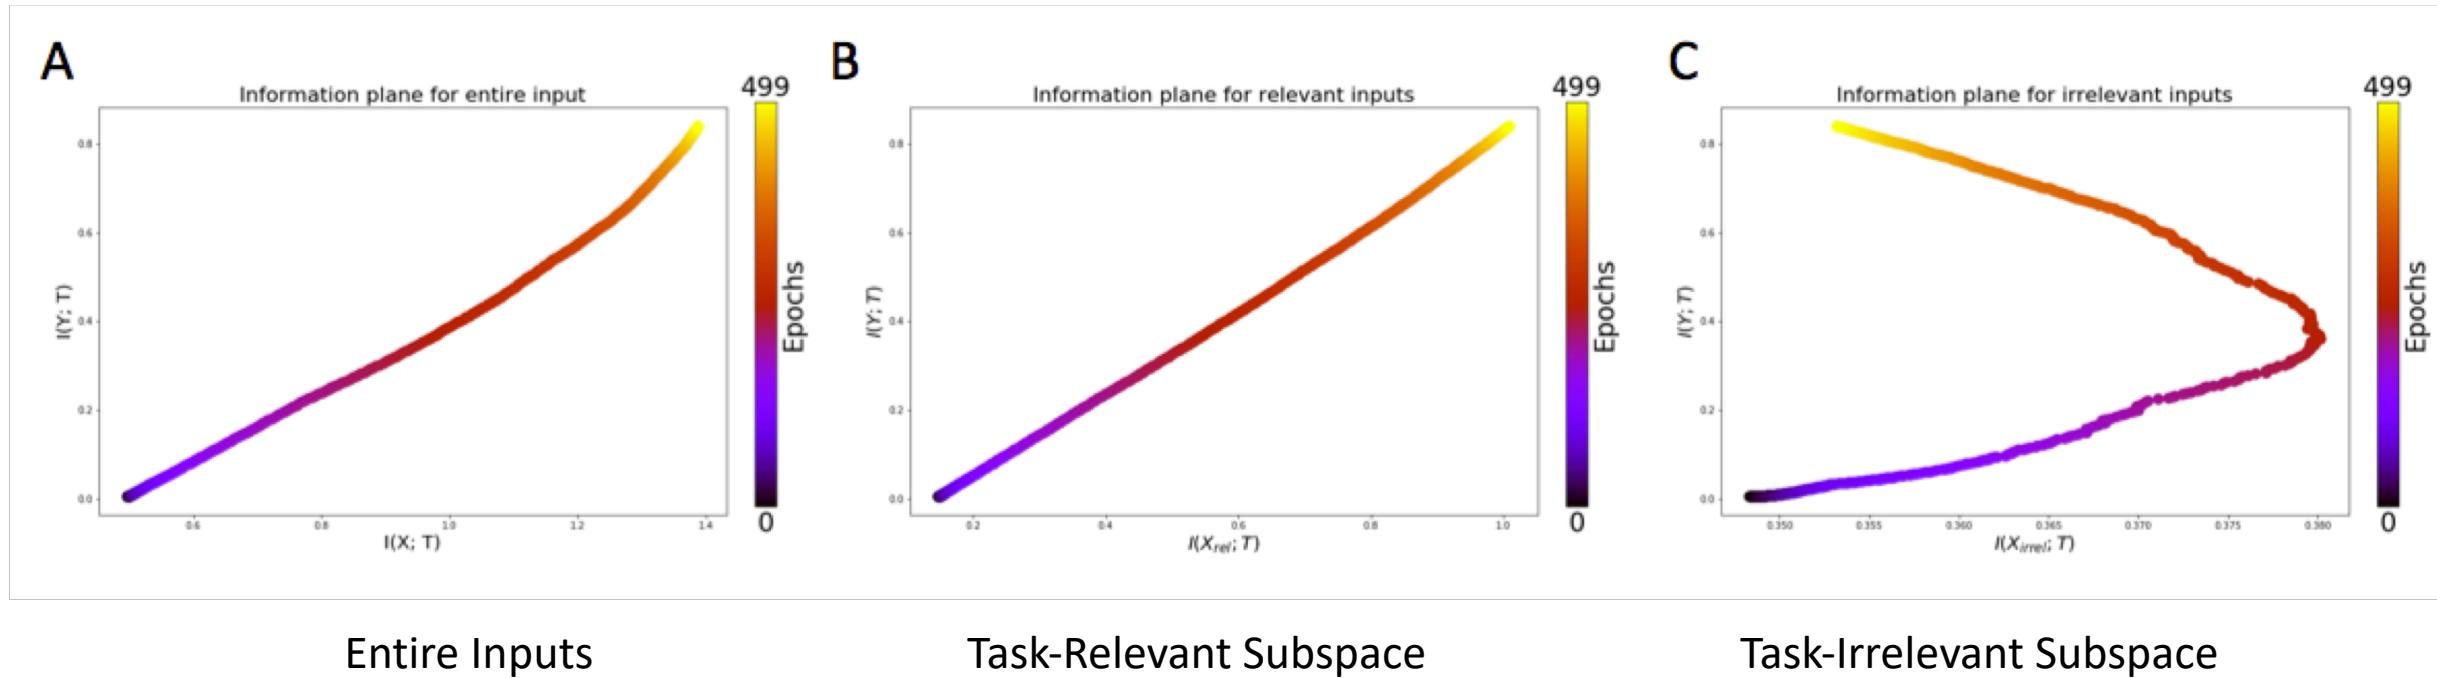
Simultaneous Fitting and Compression

- Deep Linear Network
 - Task Relevant Inputs: X_{rel}
 - Tasks Irrelevant Inputs: X_{irrel}

Signal + Noise
Only Noise

30 inputs
70 inputs

SGD
(5 samples/Batch)



Summary

- Compression dynamics
 - not general feature of deep networks
 - Critically influenced by the nonlinearities
 - Double Saturating nonlinearities lead to compression
- Compression
 - not caused by Stochasticity in training process
 - May not casually linked to generalization
- Compression
 - may still occurs in a subset of inputs
 - May not be visible from trace of information metrics
- Information Bottleneck Principle
 - May still offer important insight into deep networks
 - Could yield new training algo for inherently stochastic network, and regularization



Response from Naftali Tishby

https://www.reddit.com/r/MachineLearning/comments/79efus/r_on_the_information_bottleneck_theory_of_deep/